



## Lecture 4: Word Embedding (cont.)

Instructor: Xiang Ren  
USC CSCI 444 NLP  
2026 Spring

# Announcements

# Announcements + Logistics

- Project team finalized by end of Jan → submit your team by **Feb 2** and no more adjustment!
- HW1 is due by **Feb 4, 11:59 PM PT**
- Project proposal is due by **Feb 11, 11:59 PM PT**

# Project Pitch (cont.)

# Words as Vectors

“You shall know a word by the company it keeps.”

- Firth (1957)

# Word Meaning via Language Use

- The meaning of a word can be given by its distribution in language usage:
  - One way to define "usage": words are defined by their environments
    - Neighboring words or grammatical environments
- Intuitions: Zellig Harris (1954):
  - "oculist and eye-doctor ... occur in almost the same environments"
  - "If A and B have almost identical environments we say that they are synonyms."

A bottle of tesgüino is on the table

Everybody likes tesgüino

Tesgüino makes you drunk

We make tesgüino out of corn.

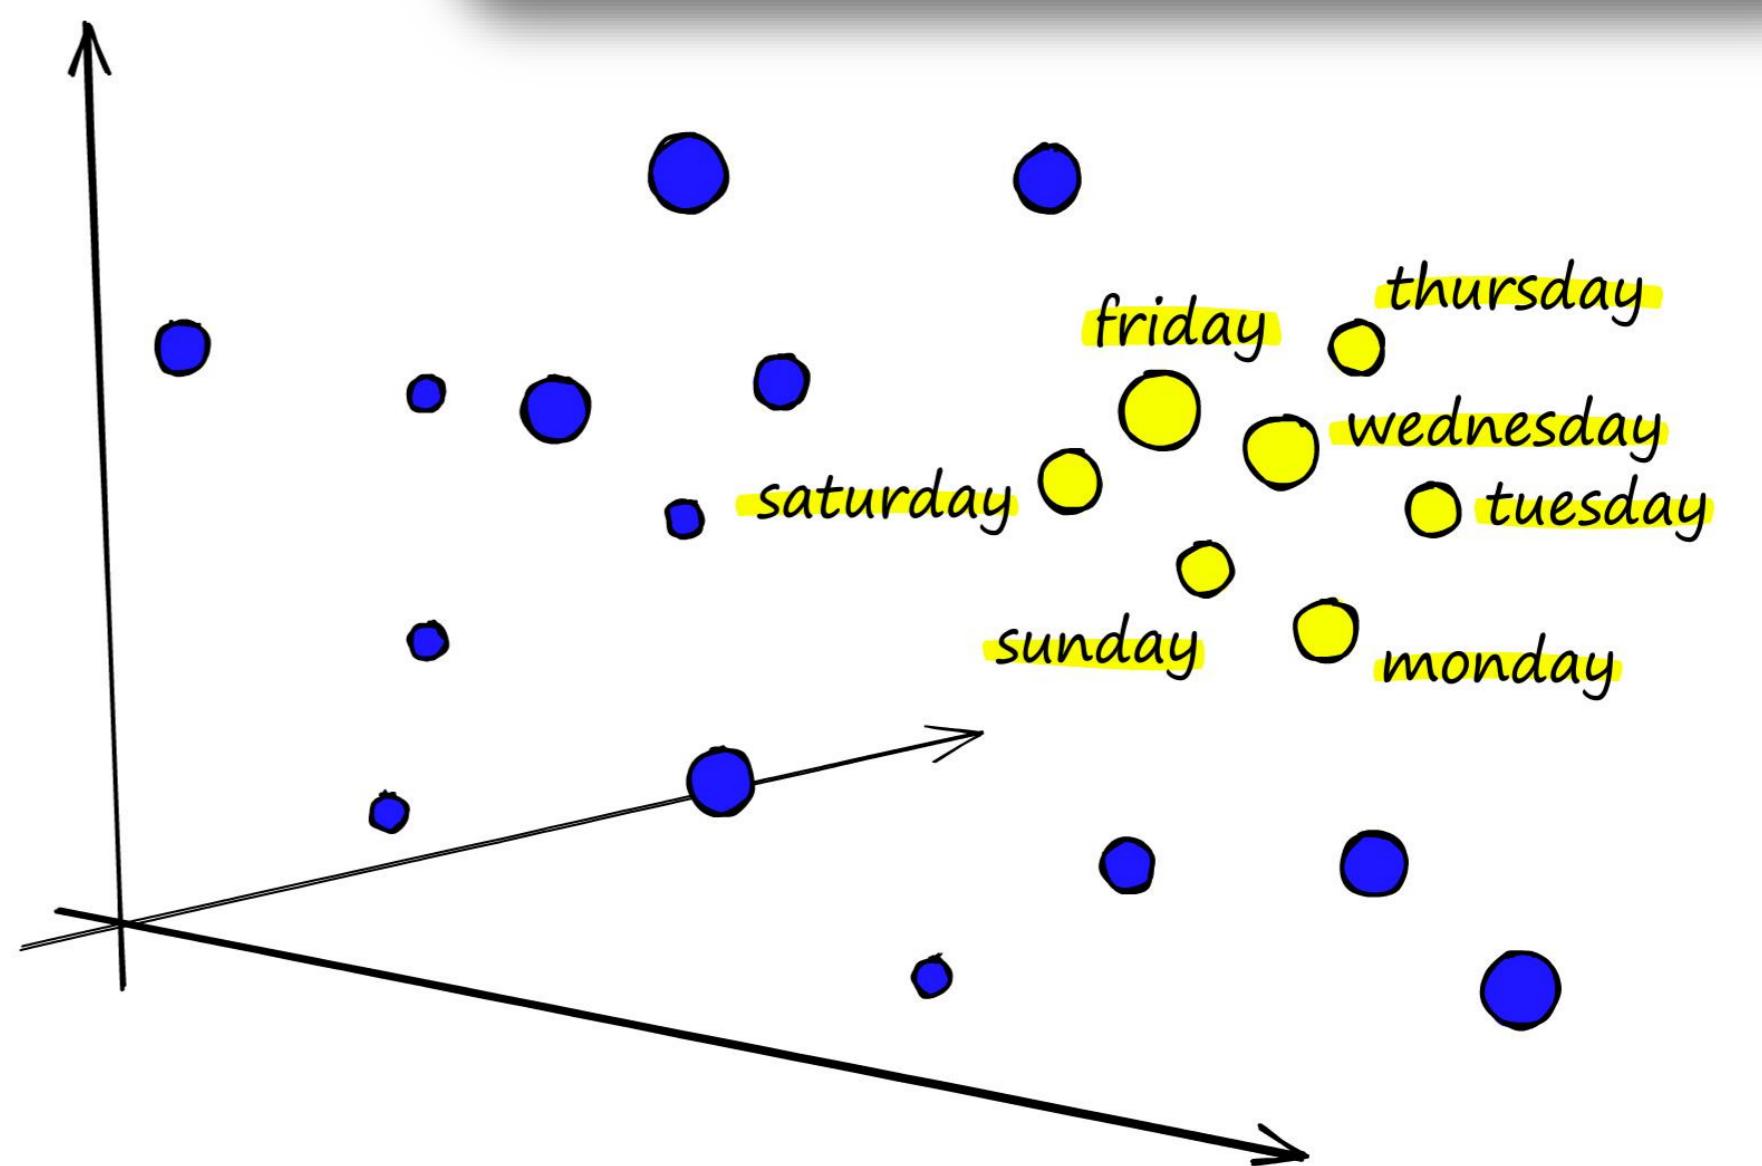


Two words are similar if they have similar word contexts

# Word Embeddings

- Represent a word as a point in a multidimensional semantic space
  - Space itself constructed from distribution of word neighbors
- Called an "embedding" because it's embedded into a space
- Fine-grained model of meaning for similarity

Vector Semantics



Every modern NLP algorithm uses embeddings as the representation of word meaning

# Intuition: Why Vectors?

- Consider generation:
  - With word strings, a feature is a word identity
    - Feature 5: ‘The previous word was “horror”’
    - Requires exact same word to be in training and test
  - With embeddings:
    - Feature is a word vector
    - ‘The previous word was vector [35,22,17...]’
    - Now in the test set we might see a similar vector [34,21,14] ... perhaps corresponding to “horrific”
    - We can generalize to similar but unseen words!!!

# Cosine Similarity for Word Similarity

Cosine similarity of two vectors

$$\cos(\vec{v}, \vec{w}) = \frac{\vec{v} \cdot \vec{w}}{\|\vec{v}\| \|\vec{w}\|}$$

$$= \frac{\sum_{i=1}^N v_i w_i}{\sqrt{\sum_{i=1}^N v_i^2} \sqrt{\sum_{i=1}^N w_i^2}}$$

Based on the definition of the dot product

between two vectors  $\vec{a}$  and  $\vec{b}$

$$\vec{v} \cdot \vec{w} = \|\vec{v}\| \|\vec{w}\| \cos \theta$$

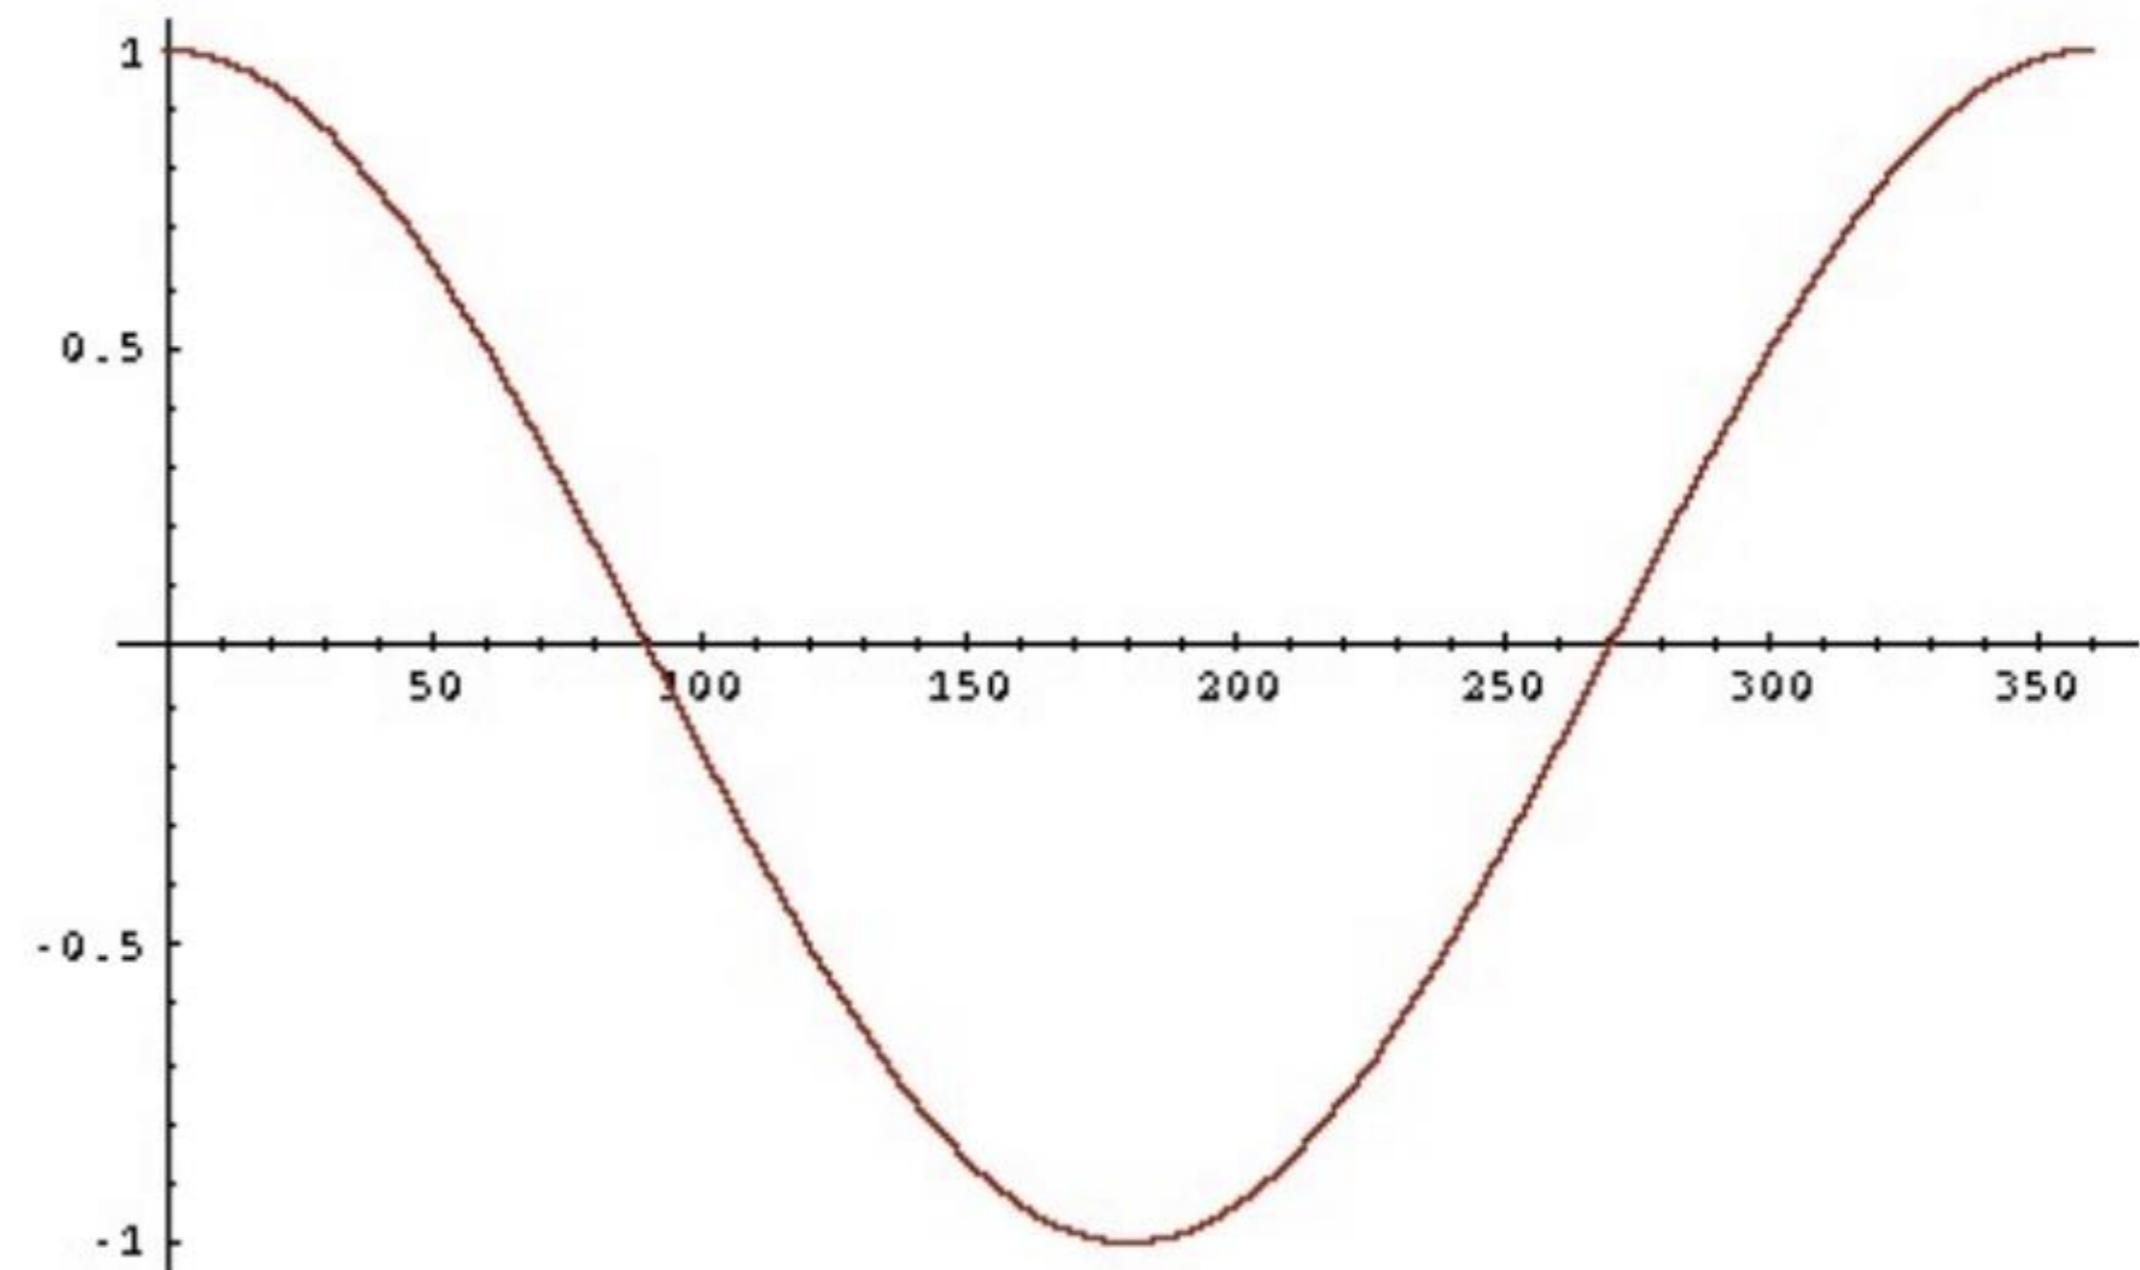
$$\cos \theta = \frac{\vec{v} \cdot \vec{w}}{\|\vec{v}\| \|\vec{w}\|}$$

# Cosine as a similarity metric

-1: vectors point in opposite directions

+1: vectors point in same directions

0: vectors are orthogonal



Greater the cosine, more similar the words

# n-grams as One-hot Vectors

**Unigram Vectors:** Represent each word as a vector of zeros with a single 1 identifying the index of the word

vocabulary

i

hate

love

the

movie

film

movie =  $<0, 0, 0, 0, 1, 0>$

film =  $<0, 0, 0, 0, 0, 1>$

One hot vector

How can we compute a vector representation such that the dot product correlates with word similarity?

Dot product is zero! These vectors are orthogonal

# Term-document matrix

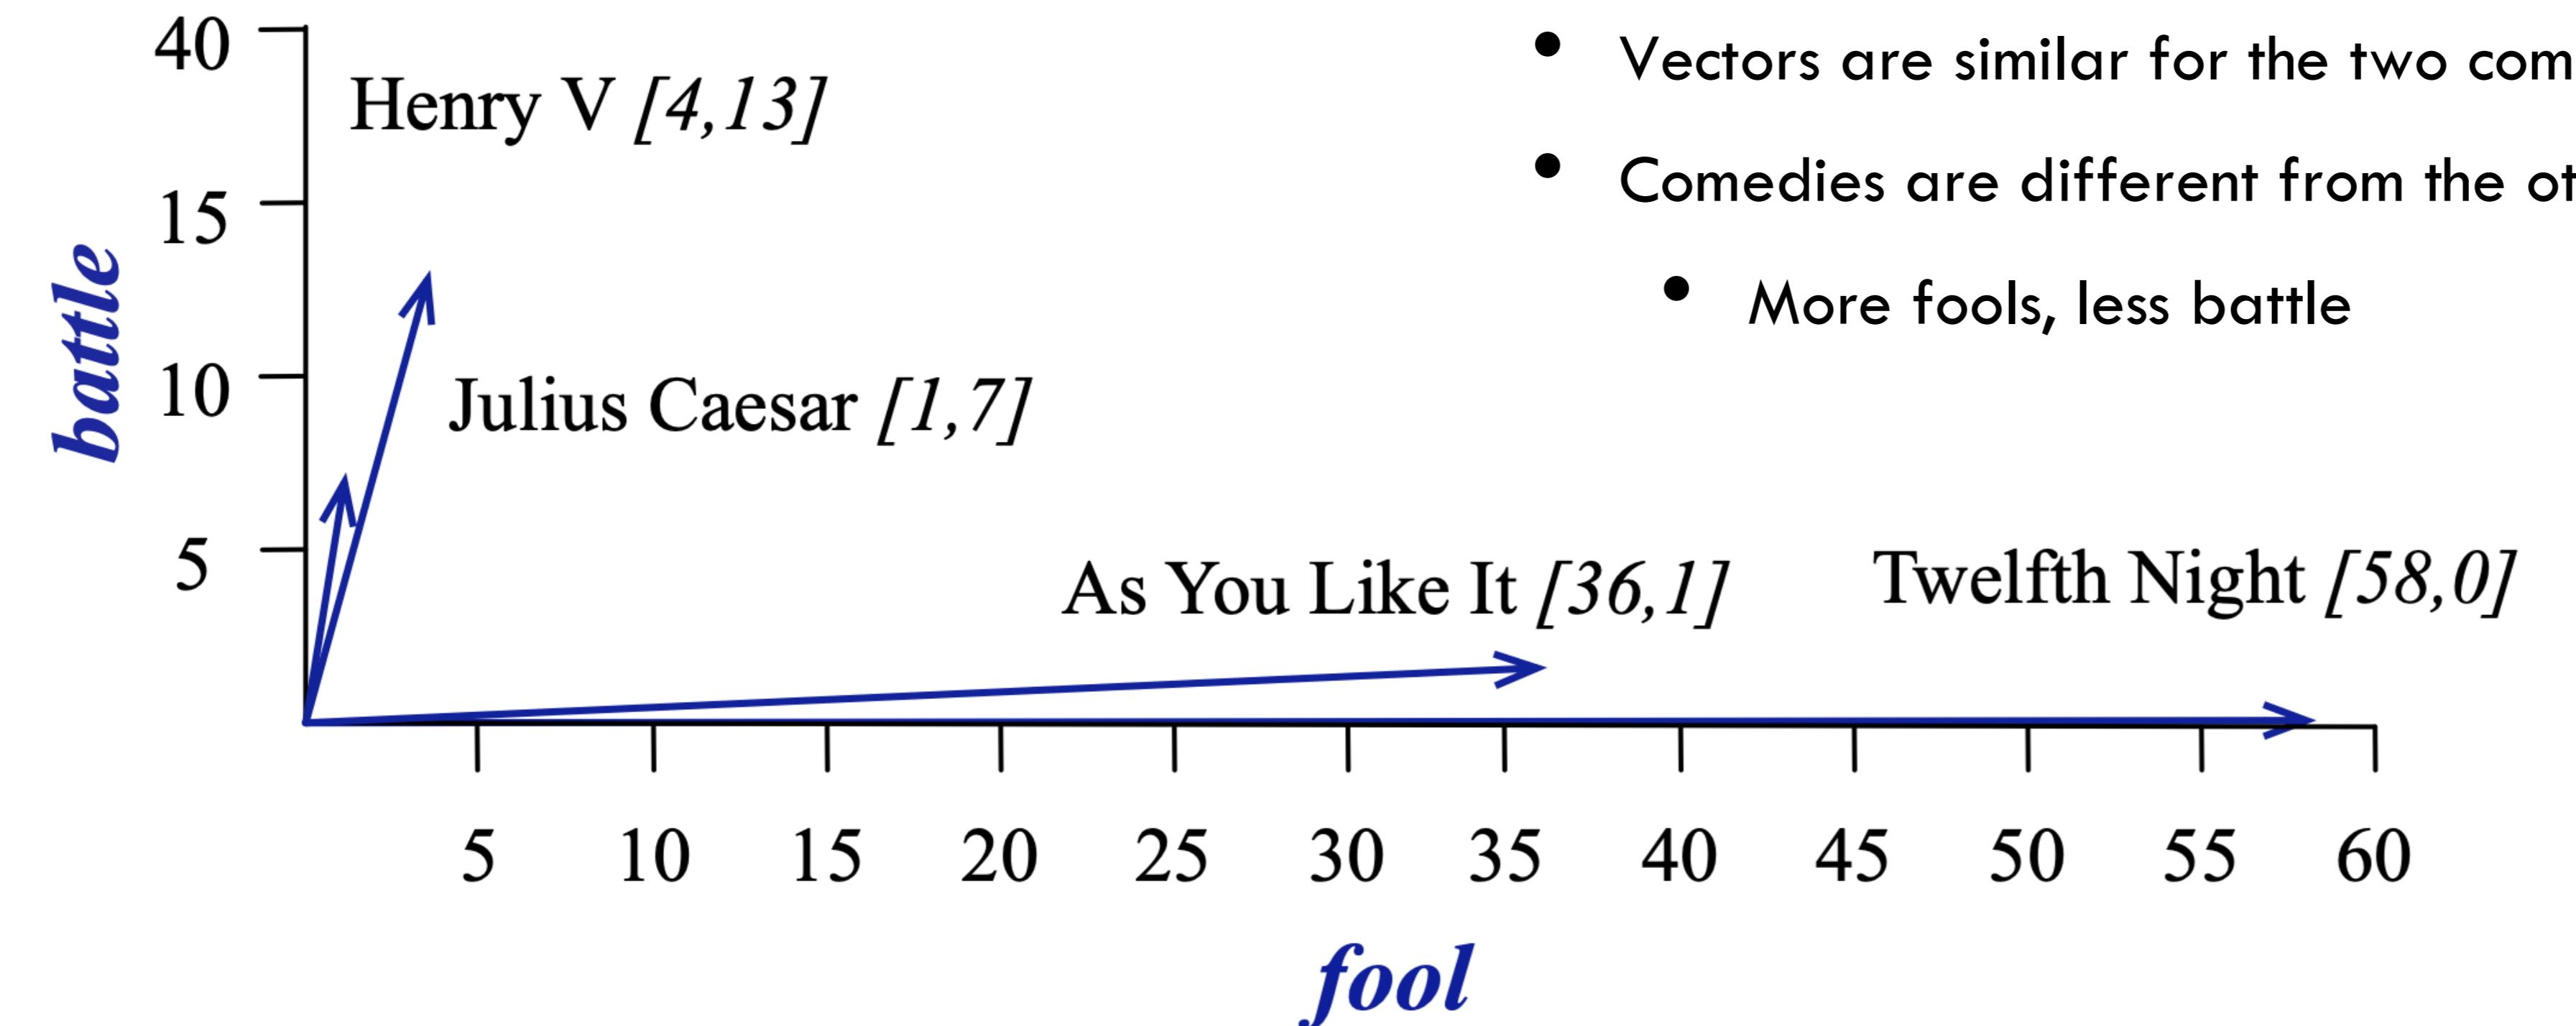
Let us consider a collection of documents and count how frequently a word (term) appears in each. A document could be a play or a Wikipedia article. In general, documents can be anything; we often call each paragraph a document!

Each **document** is represented by a vector of words

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	0	7	13
good	114	80	62	89
fool	36	58	1	4
wit	20	15	2	3

# Visualizing document vectors

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	0	7	13
good	14	80	62	89
fool	36	58	1	4
wit	20	15	2	3



# Words as vectors in a co-occurrence matrix

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	0	7	13
good	114	80	62	89
fool	36	58	1	4
wit	20	15	2	3

“Battle” is the kind of word that appears in Julius Caesar and Henry V

“Fool” is the kind of word that appears in As You Like It and Twelfth Night

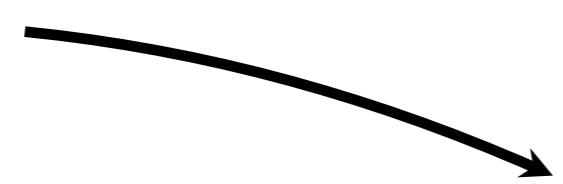
# Word-word co-occurrence matrix

Two words are similar in meaning if their context vectors are similar

Context  
Window

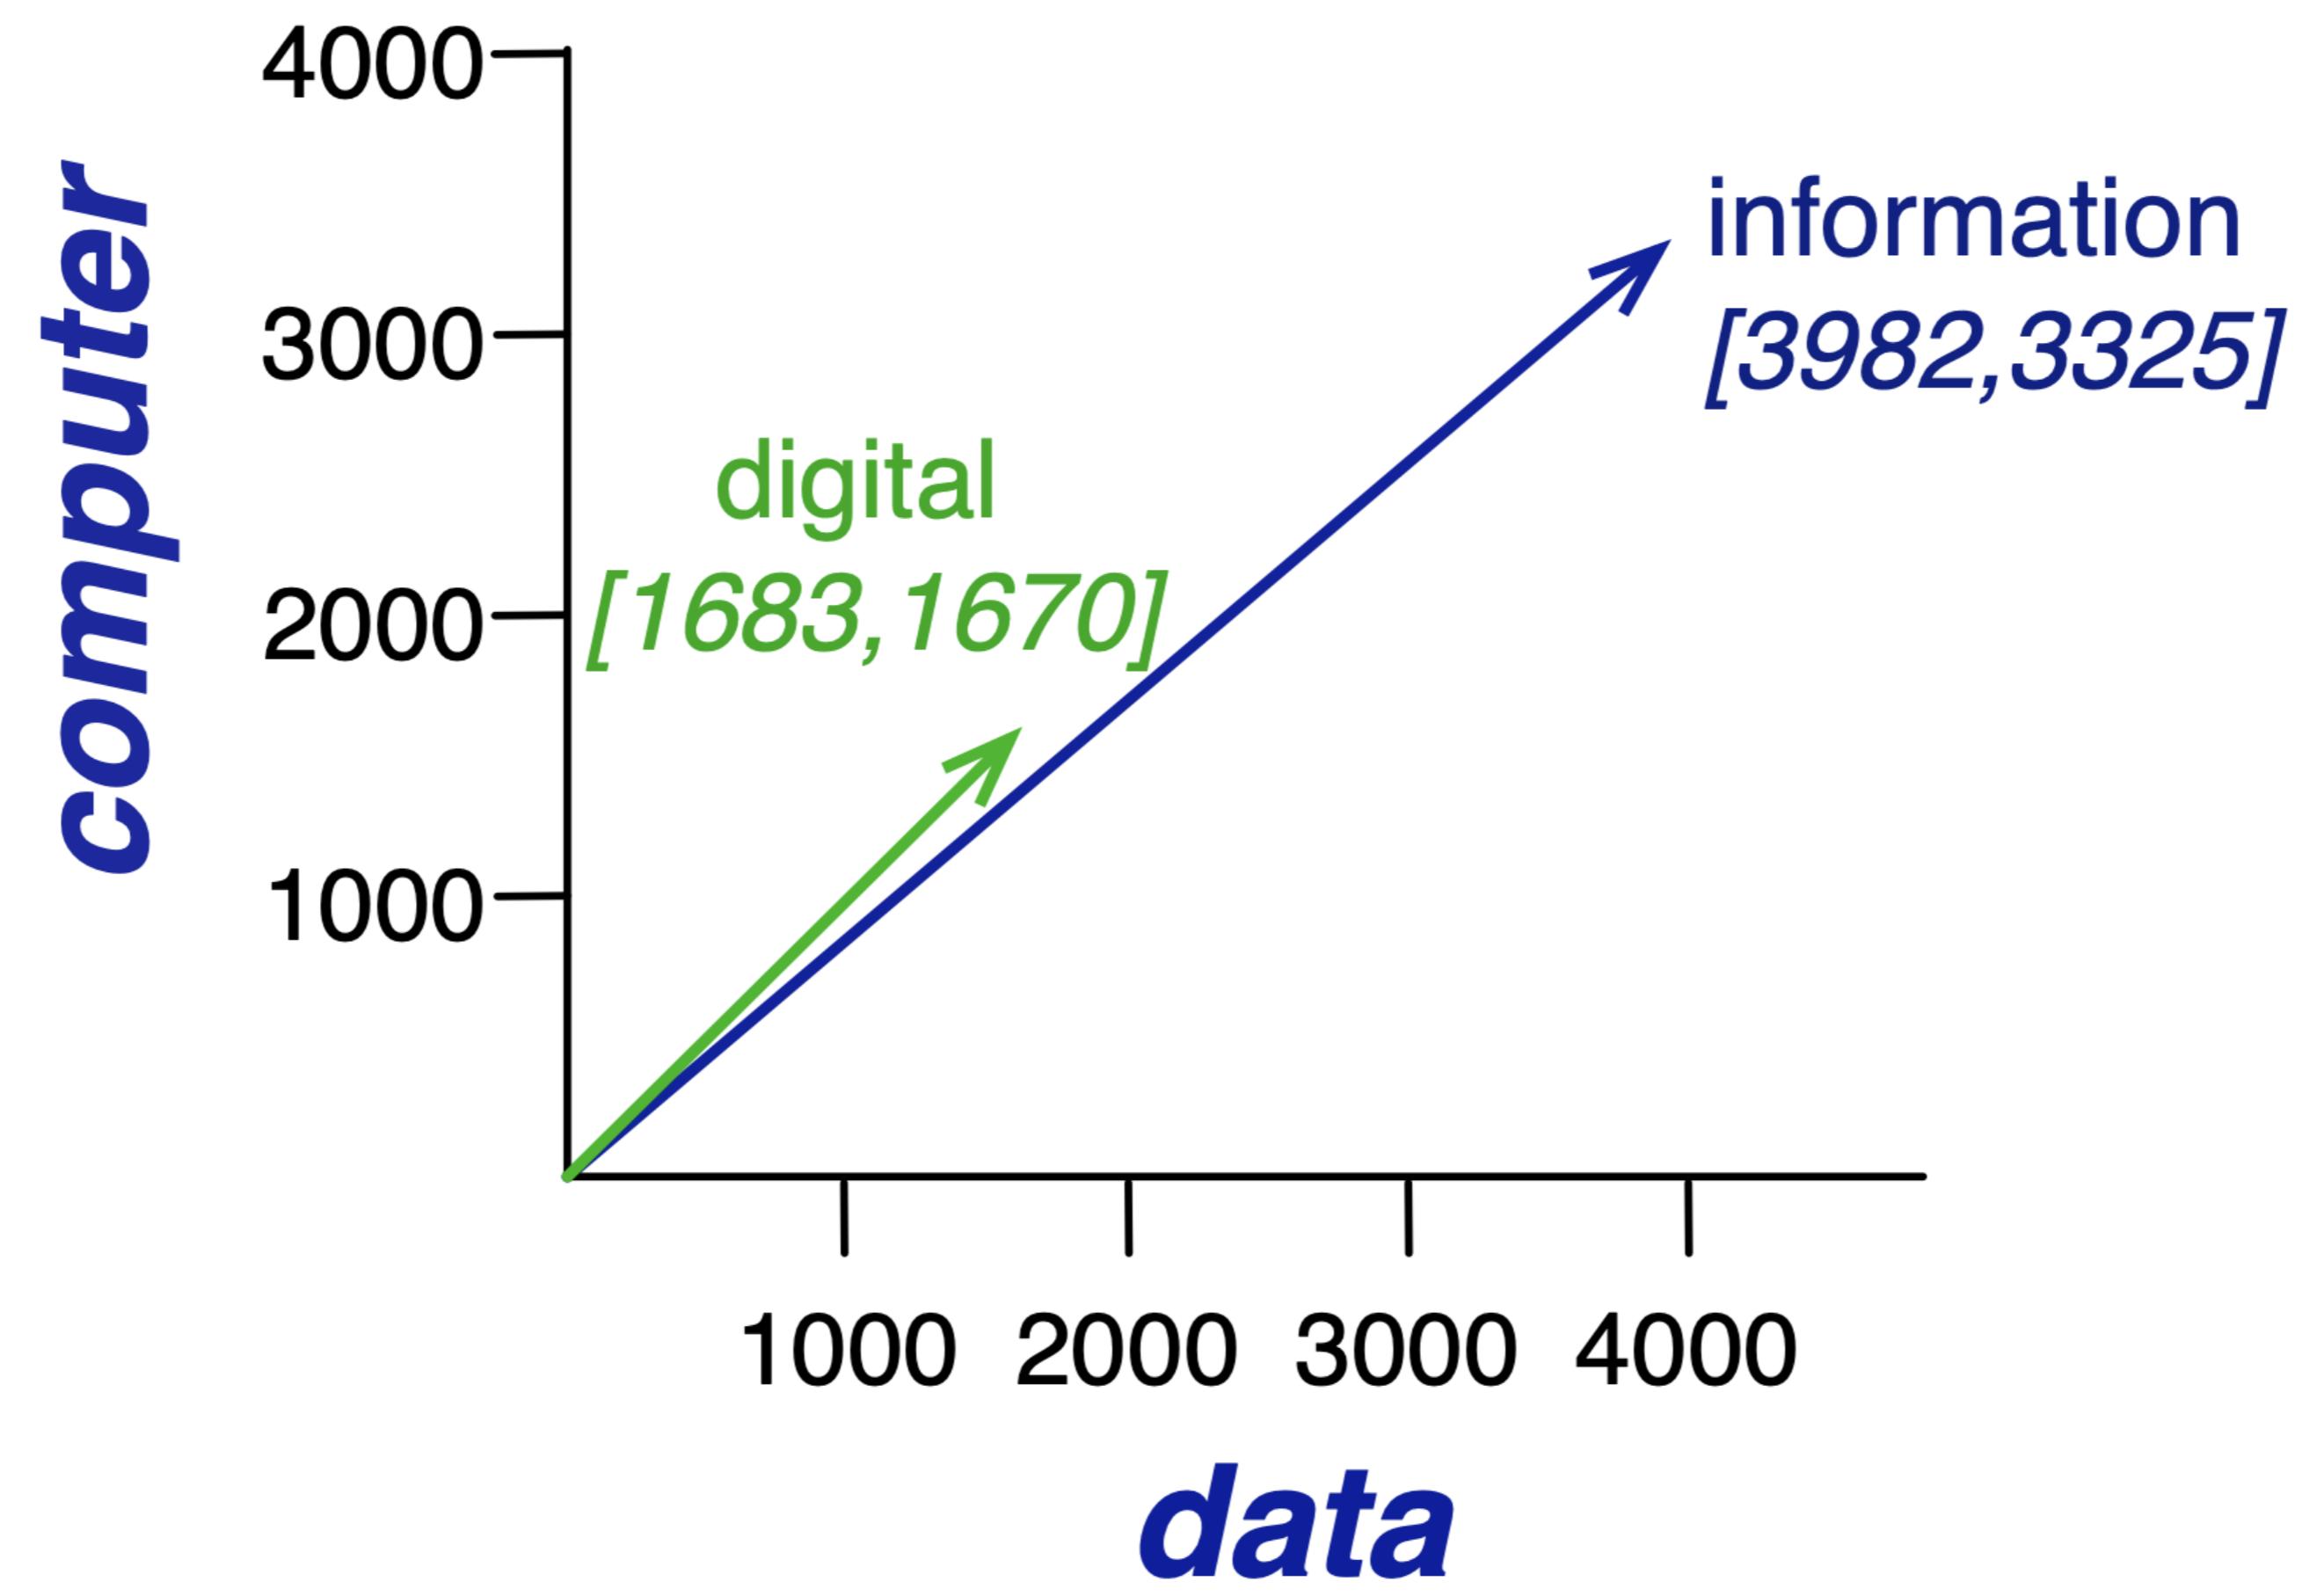
is traditionally followed by **cherry** pie, a traditional dessert  
often mixed, such as **strawberry** rhubarb pie. Apple pie  
computer peripherals and personal **digital** assistants. These devices usually  
a computer. This includes **information** available on the internet

Words, not documents



	aardvark	...	computer	data	result	pie	sugar	...
<b>cherry</b>	0	...	2	8	9	442	25	...
<b>strawberry</b>	0	...	0	0	1	60	19	...
<b>digital</b>	0	...	1670	1683	85	5	4	...
<b>information</b>	0	...	3325	3982	378	5	13	...

	aardvark	...	computer	data	result	pie	sugar	...
cherry	0	...	2	8	9	442	25	...
strawberry	0	...	0	0	1	60	19	...
digital	0	...	1670	1683	85	5	4	...
information	0	...	3325	3982	378	5	13	...



# Cosine Similarity

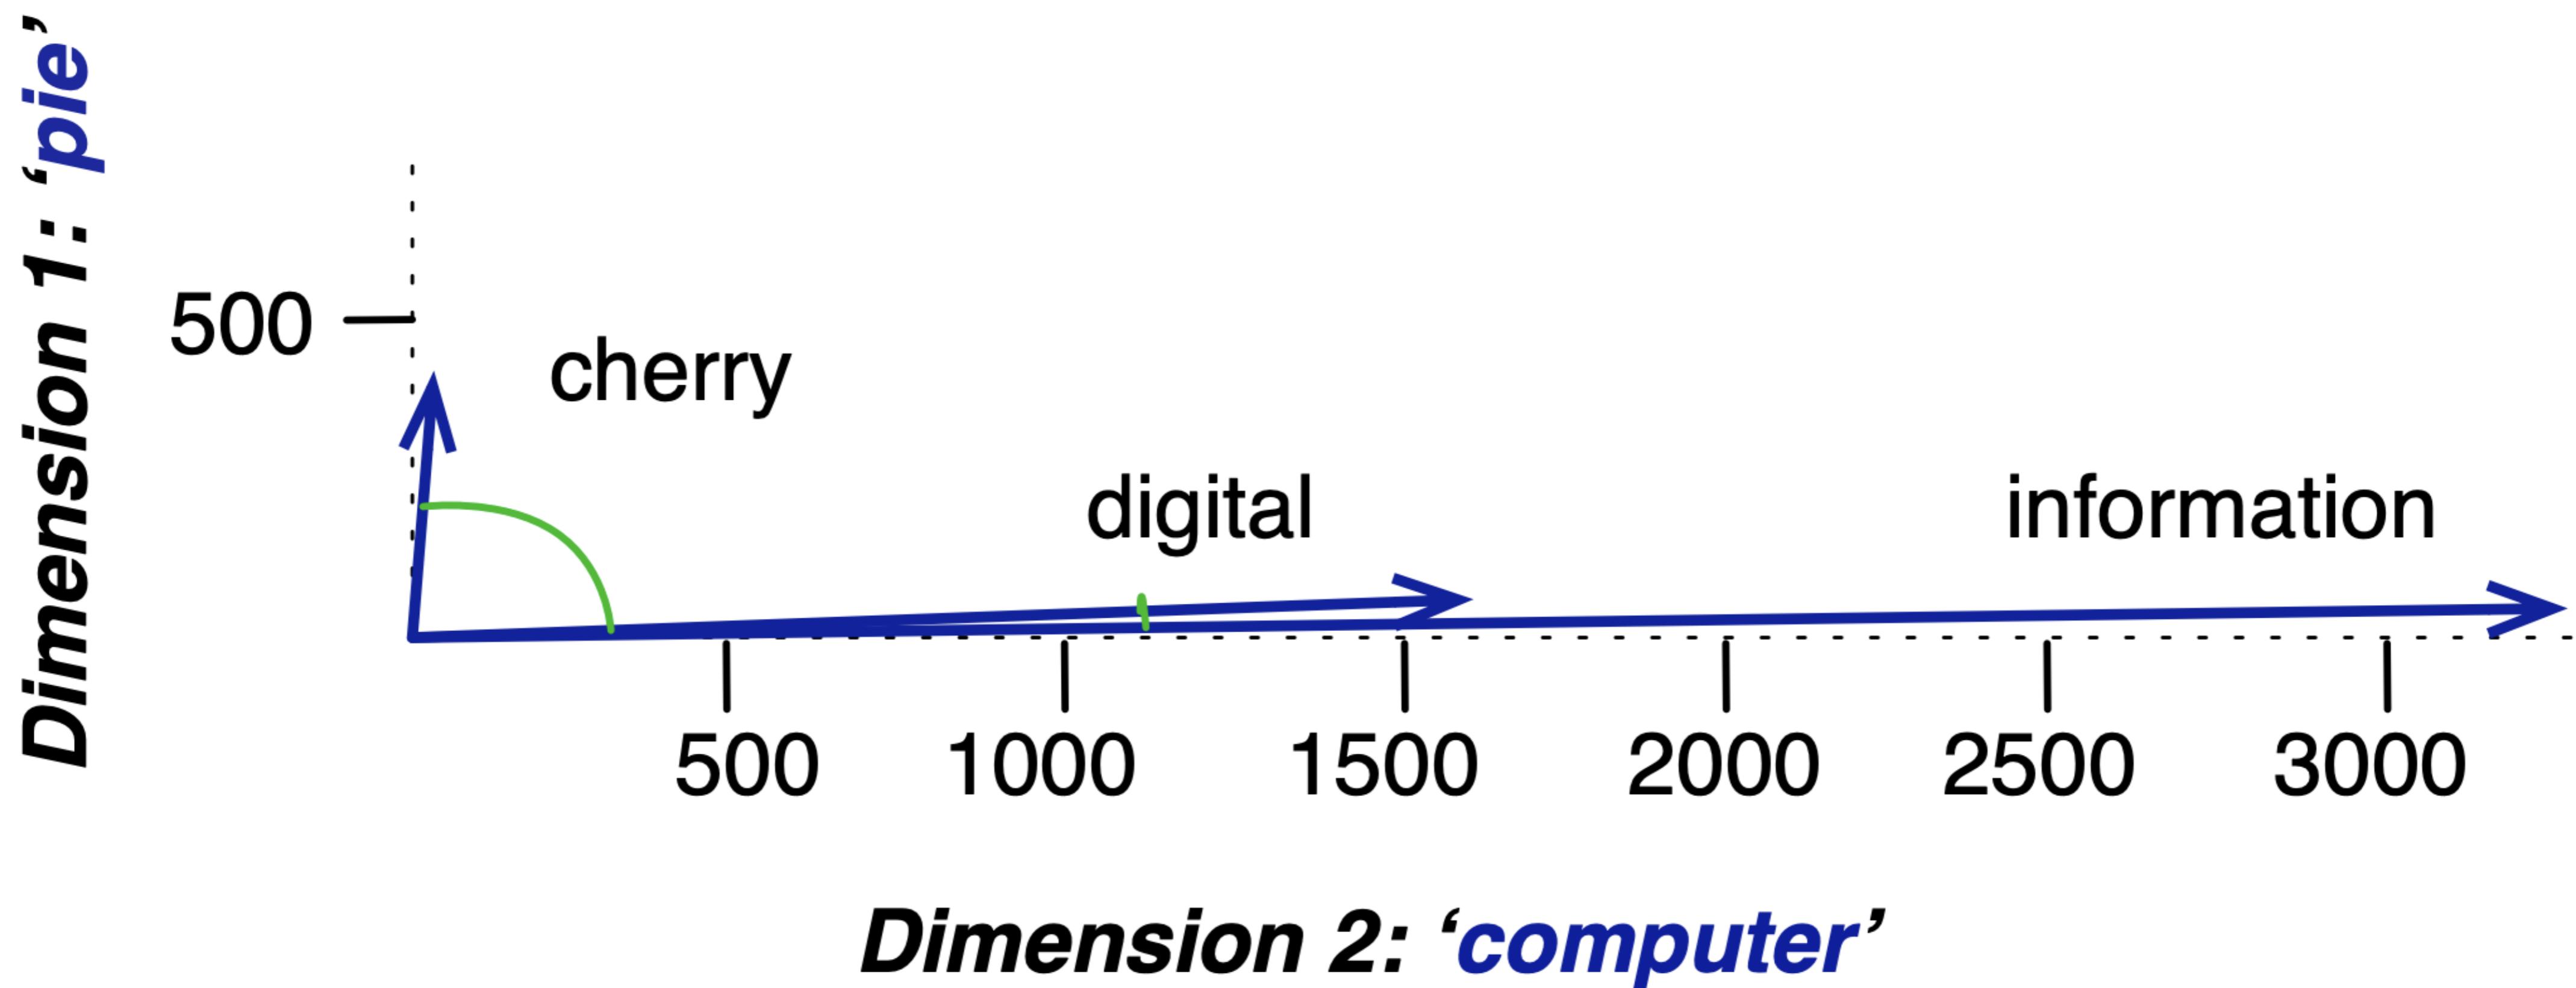
$$\begin{aligned} \cos(\vec{v}, \vec{w}) &= \frac{\vec{v} \cdot \vec{w}}{|\vec{v}| |\vec{w}|} \\ &= \frac{\sum_{i=1}^N v_i w_i}{\sqrt{\sum_{i=1}^N v_i^2} \sqrt{\sum_{i=1}^N w_i^2}} \end{aligned}$$

	<b>pie</b>	<b>data</b>	<b>computer</b>
cherry	442	8	2
digital	5	1683	1670
information	5	3982	3325

$$\cos(\text{cherry}, \text{information}) = \frac{442 * 5 + 8 * 3982 + 2 * 3325}{\sqrt{442^2 + 8^2 + 2^2} \sqrt{5^2 + 3982^2 + 3325^2}} = .017$$

$$\cos(\text{digital}, \text{information}) = \frac{5 * 5 + 1683 * 3982 + 1670 * 3325}{\sqrt{5^2 + 1683^2 + 1670^2} \sqrt{5^2 + 3982^2 + 3325^2}} = .996$$

# Visualizing cosines



# Raw frequencies though...

- ...are a bad representation!
- The co-occurrence matrices we have seen represent each cell by word frequencies
- Frequency is clearly useful; if sugar appears a lot near apricot, that's useful information
- But overly frequent words like the, it, or they are not very informative about the context
- It's a paradox! How can we balance these two conflicting constraints?

	pie	data	computer
cherry	442	8	2
digital	5	1683	1670
information	5	3982	3325

Need some form of weighting!

# Two different kinds of weighting

## **tf-idf: Term Frequency - Inverse Document Frequency**

- Downweighting words like “the” or “if”
- Term-document matrices

## **PMI: Pointwise Mutual Information**

- Considers the probability of words like “good” and “great” co-occurring
- Word co-occurrence matrices

# Term Frequency

Term Frequency: frequency counting (usually log transformed)

$$tf_{t,d} = \begin{cases} 1 + \log(count(t, d)), & \text{if } count(t, d) > 0 \\ 0, & \text{otherwise} \end{cases}$$

$count(t, d)$  = # occurrences of word  $t$  in document  $d$

# Inverse Document Frequency

- Document Frequency:  $df_t$  is the number of documents  $t$  occurs in.

	Collection Frequency	Document Frequency
Romeo	113	1
action	113	31

- NOT collection frequency: total count across all documents
- "Romeo" is very distinctive for one Shakespeare play
- Inverse Document Frequency:  $idf_t$

$$idf_t = \log_{10} \left( \frac{N}{df_t} \right)$$

$N$  = total number of documents in the collection

Word	df	idf
Romeo	1	1.57
salad	2	1.27
Falstaff	4	0.967
forest	12	0.489
battle	21	0.246
wit	34	0.037
fool	36	0.012
good	37	0
sweet	37	0



# tf-idf

Final tf-idf weighted value for a word

$$tf_{t,d} \times idf_{t,d}$$

	<b>As You Like It</b>	<b>Twelfth Night</b>	<b>Julius Caesar</b>	<b>Henry V</b>
<b>battle</b>	1	0	7	13
<b>good</b>	114	80	62	89
<b>fool</b>	36	58	1	4
<b>wit</b>	20	15	2	3

Raw Counts

tf-idf Weighted Counts

	<b>As You Like It</b>	<b>Twelfth Night</b>	<b>Julius Caesar</b>	<b>Henry V</b>
<b>battle</b>	0.074	0	0.22	0.28
<b>good</b>	0	0	0	0
<b>fool</b>	0.019	0.021	0.0036	0.0083
<b>wit</b>	0.049	0.044	0.018	0.022

# Pointwise Mutual Information (PMI)

- PMI between two words:
  - Do words  $x$  and  $y$  co-occur more than if they were independent?
- PMI ranges from  $-\infty$  to  $+\infty$ 
  - Negative values are problematic: words are co-occurring less than we expect by chance
  - Only reliable under an enormous corpora
    - Imagine  $w_1$  and  $w_2$  whose probability is each  $10^{-6}$
    - Hard to be sure  $P(w_1, w_2)$  is significantly different than  $10^{-12}$
    - So we just replace negative PMI values by 0
- Positive PMI

$$PMI(w_1, w_2) = \log \frac{P(w_1, w_2)}{P(w_1)P(w_2)}$$

$$PPMI(w_1, w_2) = \max\left(0, \log \frac{P(w_1, w_2)}{P(w_1)P(w_2)}\right)$$

# Computing PPMI on a term-context matrix

	Context $C$					
	computer	data	result	pie	sugar	count(w)
Term $W$	cherry	2	8	9	442	25
	strawberry	0	0	1	60	19
	digital	1670	1683	85	5	4
	information	3325	3982	378	5	13
	count(context)	4997	5673	473	512	61
						11716

Frequency  $f(w_i, c_j)$  or  $f_{ij}$  is the # times  $w_i$  occurs in context  $c_j$

$$P_{ij} = \frac{f_{ij}}{\sum_{i=1}^W \sum_{j=1}^C f_{i,j}} \quad P_i = \sum_{j=1}^C P_{ij} \quad P_j = \sum_{i=1}^W P_{ij}$$

$$PPMI(w_i, c_j) = PPMI_{i,j} = \max\left(0, \log \frac{P_{ij}}{P_i P_j}\right)$$

	<b>computer</b>	<b>data</b>	<b>result</b>	<b>pie</b>	<b>sugar</b>	<b>count(w)</b>
<b>cherry</b>	2	8	9	442	25	486
<b>strawberry</b>	0	0	1	60	19	80
<b>digital</b>	1670	1683	85	5	4	3447
<b>information</b>	3325	3982	378	5	13	7703
<b>count(context)</b>	4997	5673	473	512	61	11716

$$P_{ij} = \frac{f_{ij}}{\sum_{i=1}^W \sum_{j=1}^C f_{i,j}}$$

$$P_i = \sum_{j=1}^C P_{ij}$$

$$P_j = \sum_{i=1}^W P_{ij}$$

	<b>p(w,context)</b>					<b>p(w)</b>
	<b>computer</b>	<b>data</b>	<b>result</b>	<b>pie</b>	<b>sugar</b>	<b>p(w)</b>
<b>cherry</b>	0.0002	0.0007	0.0008	0.0377	0.0021	0.0415
<b>strawberry</b>	0.0000	0.0000	0.0001	0.0051	0.0016	0.0068
<b>digital</b>	0.1425	0.1436	0.0073	0.0004	0.0003	0.2942
<b>information</b>	0.2838	0.3399	0.0323	0.0004	0.0011	0.6575
<b>p(context)</b>	0.4265	0.4842	0.0404	0.0437	0.0052	

$$p(w=\text{information}, c=\text{data}) = 3982/111716 = .3399$$

$$p(w=\text{information}) = 7703/111716 = .6575$$

$$p(c=\text{data}) = 5673/111716 = .4842$$

	<b>p(w,context)</b>					<b>p(w)</b>
	<b>computer</b>	<b>data</b>	<b>result</b>	<b>pie</b>	<b>sugar</b>	<b>p(w)</b>
<b>cherry</b>	0.0002	0.0007	0.0008	0.0377	0.0021	0.0415
<b>strawberry</b>	0.0000	0.0000	0.0001	0.0051	0.0016	0.0068
<b>digital</b>	0.1425	0.1436	0.0073	0.0004	0.0003	0.2942
<b>information</b>	0.2838	0.3399	0.0323	0.0004	0.0011	0.6575
<b>p(context)</b>	0.4265	0.4842	0.0404	0.0437	0.0052	

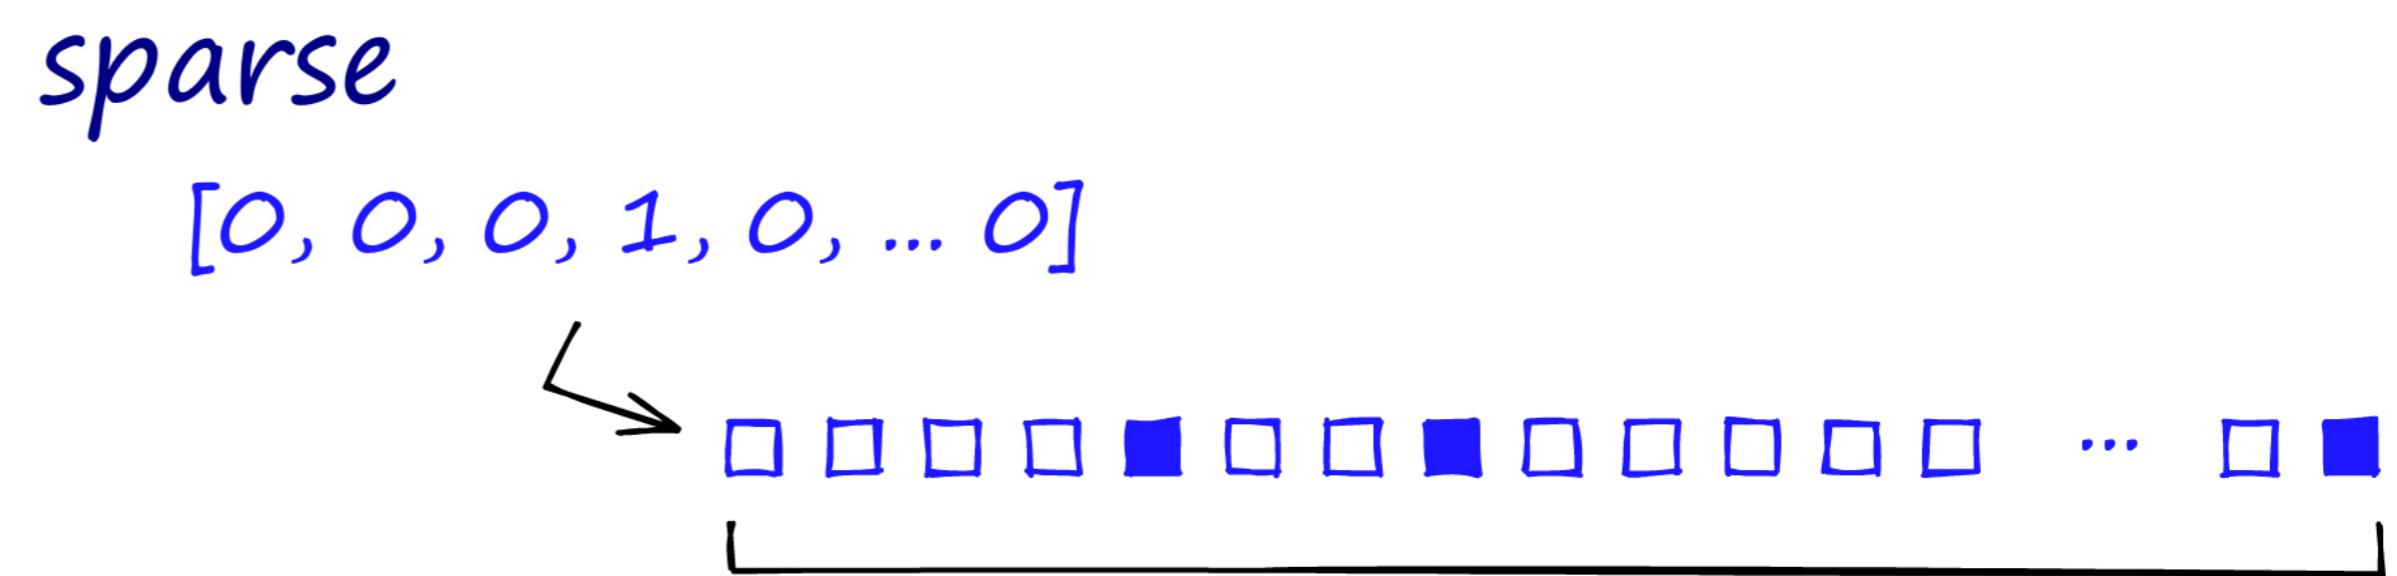
	<b>computer</b>	<b>data</b>	<b>result</b>	<b>pie</b>	<b>sugar</b>
<b>cherry</b>	0	0	0	4.38	3.30
<b>strawberry</b>	0	0	0	4.10	5.51
<b>digital</b>	0.18	0.01	0	0	0
<b>information</b>	0.02	0.09	0.28	0	0

$$\text{pmi}(\text{information}, \text{data}) = \log_2 (.3399 / (.6575 * .4842)) = .0944$$

$$PPMI_{i,j} = \max\left(0, \log \frac{P_{ij}}{P_i P_j}\right)$$

# The problem...

- tf-idf (or PMI) vectors are
  - long (length  $|V| = 20,000$  to  $50,000$ )
  - sparse (most elements are zero)



- Alternative: learn vectors which are
  - short (length 50-1000)
  - dense (most elements are non-zero)

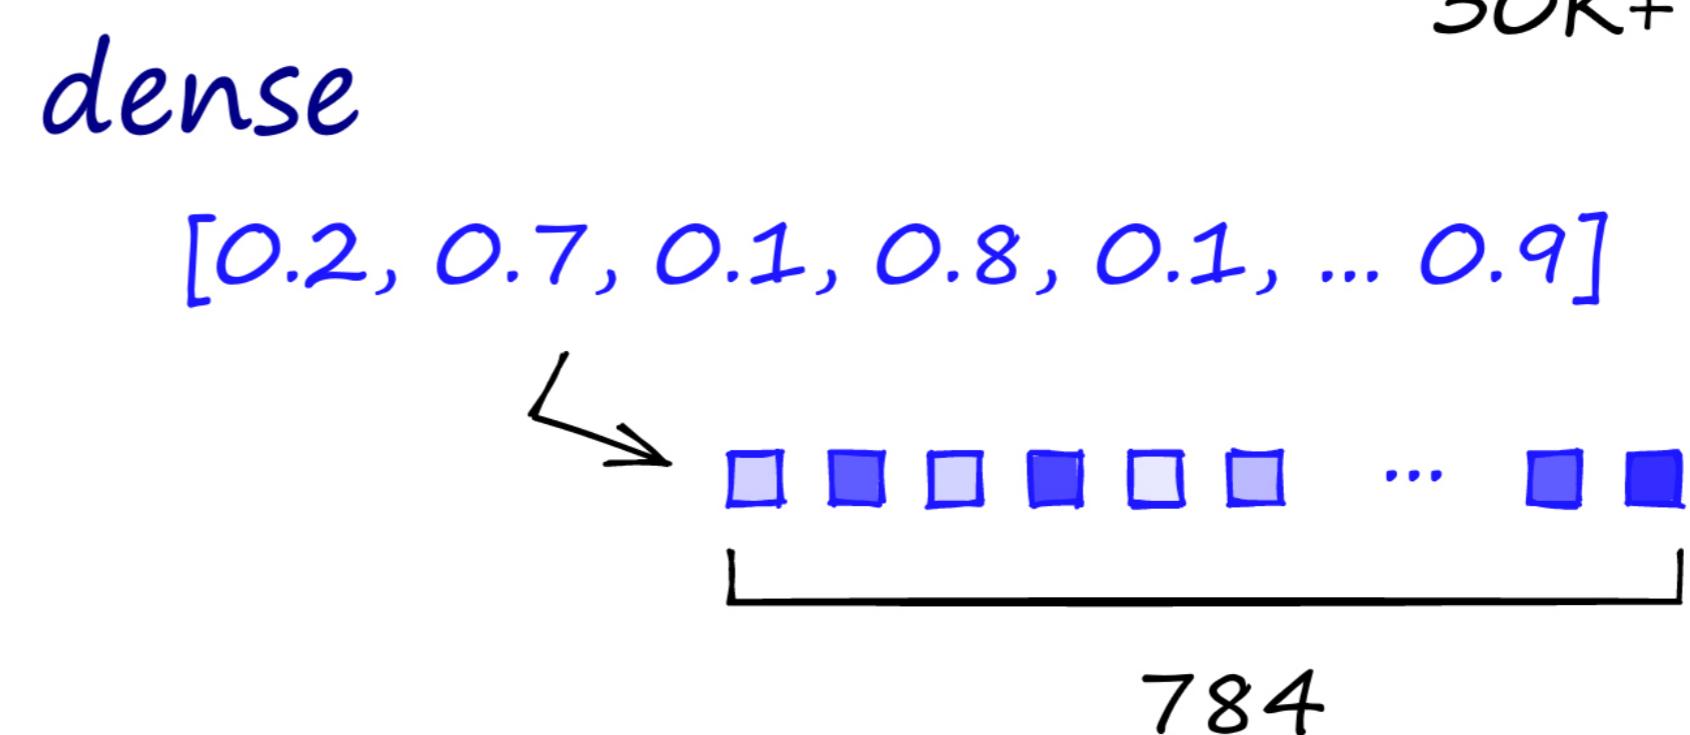


Image Credit: Pinecone

# Sparse vs. Dense Vectors

- Why dense vectors?
  - Memory efficiency is not a problem for sparse vectors...
  - Short vectors may be easier to use as features in machine learning (fewer weights to tune)
  - Dense vectors may generalize better than explicit counts
  - Dense vectors may do better at capturing synonymy:
    - car and automobile are synonyms; but are distinct dimensions
    - a word with car as a neighbor and a word with automobile as a neighbor should be similar, but aren't
  - In practice, they work better



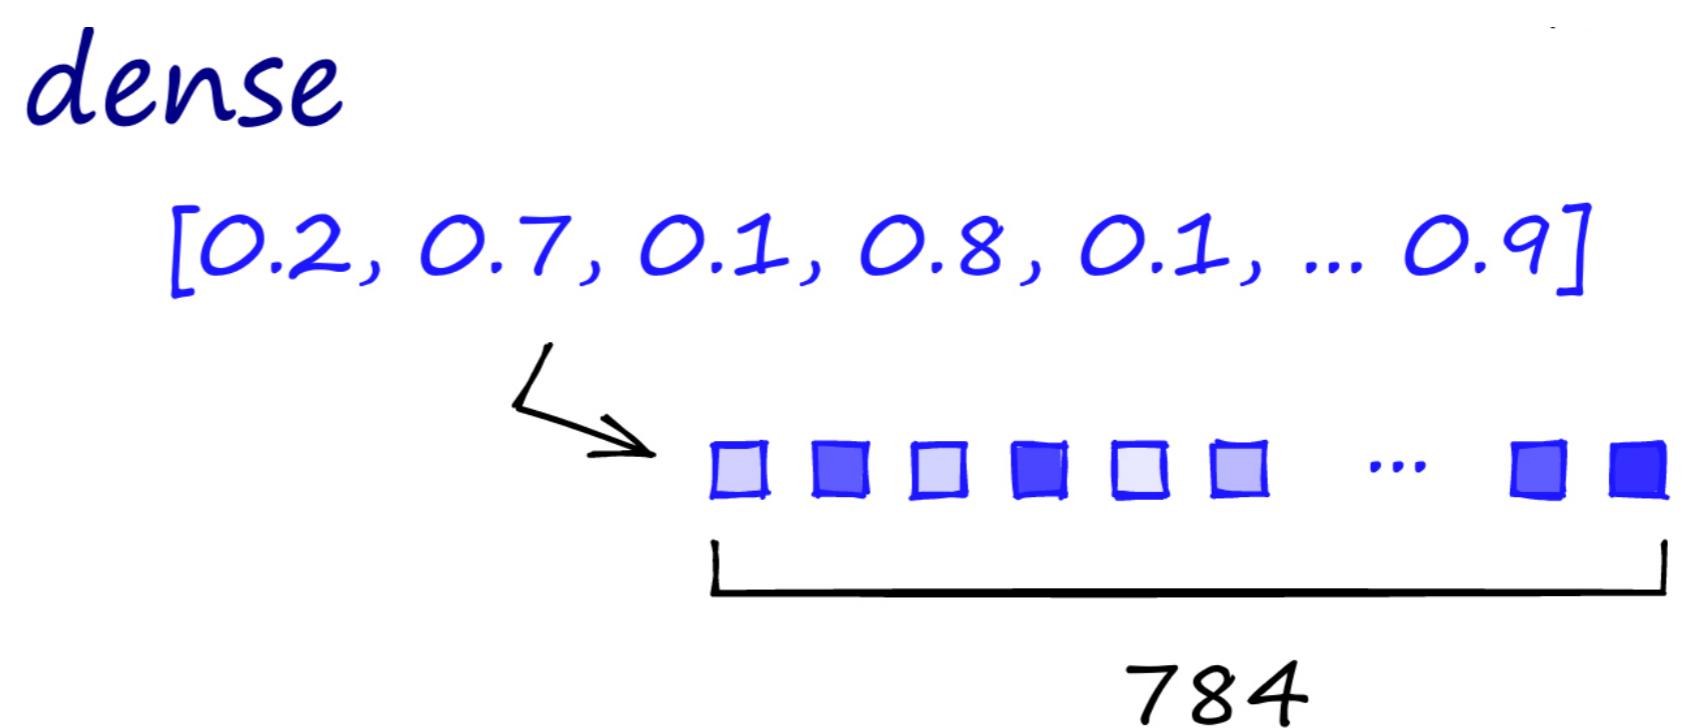
# How to obtain dense vectors?

“Neural Language Model”-inspired models

- Word2vec (skipgram, CBOW), GloVe

Singular Value Decomposition (SVD)

- Special case: Latent Semantic Analysis (LSA)



Alternative to these "static embeddings":

- Contextual Embeddings (ELMo, BERT)
- Compute distinct embeddings for a word in its context
- Separate embeddings for each token of a word

# word2vec

# word2vec

Mikolov et al., ICLR 2013. Efficient estimation of word representations in vector space.

Mikolov et al., NeurIPS 2013. Distributed representations of words and phrases and their compositionality.

- Short, dense vector or embedding
- Static embeddings
  - One embedding per word type
  - Does not change with context change
- Two algorithms for computing:
  - Skip-Gram with Negative Sampling or SGNS
  - CBOW or continuous bag of words
  - But we will study a slightly different version...
- Efficient training
- Easily available to download and plug in

Polysemy?



# word2vec : Intuition

is traditionally followed by **cherry** pie, a traditional dessert  
often mixed, such as **strawberry** rhubarb pie. Apple pie  
computer peripherals and personal **digital** assistants. These devices usually  
a computer. This includes **information** available on the internet

Instead of counting how often each word  $W$  occurs near another, e.g. “cherry”

- Train a classifier on a binary prediction task:
  - Is  $W$  likely to show up near “cherry”?
- We don't actually care about this task!!!
- But we'll take the learned **classifier weights** as the word embeddings

$x?$   $y?$

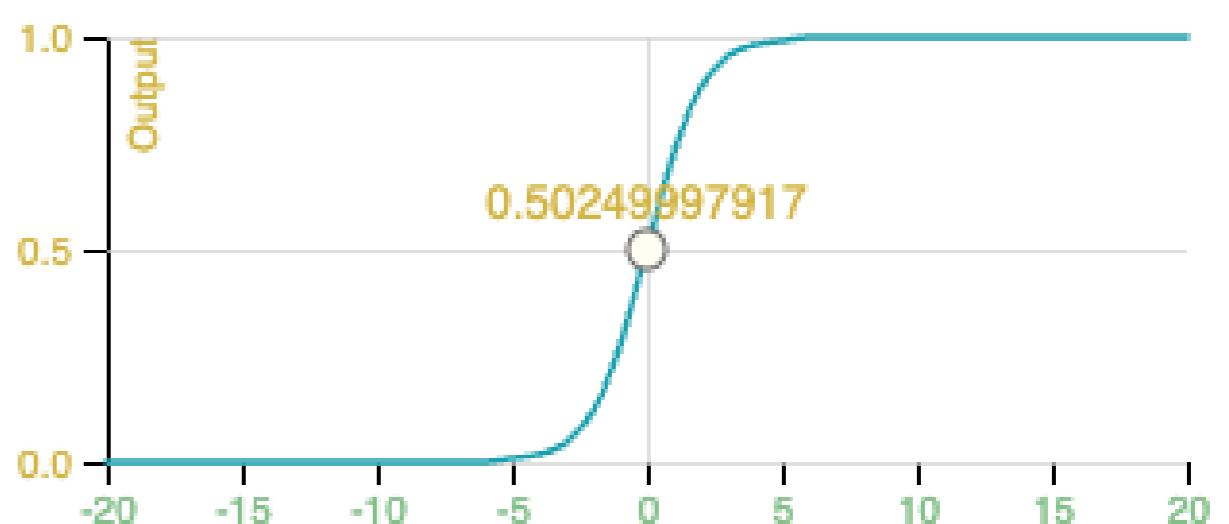


Word embedding itself is the learned parameter!

# Binary Text Classification

- Goal: Given an input, predict label or class from a discrete set
  - e.g. Predict the sentiment (positive or negative) for a sentence
- Input:  $x$  represented by feature vector of size  $d$ , given by  $\mathbf{x} \in \mathbb{R}^d$
- Output:  $y \in \{0,1\}$  for binary classification
- Suffices to learn conditional probabilities
  - Parameterized by  $\theta \in \mathbb{R}^d$
- Could estimate by cooccurrence counts, but a single feature
  - Better option: dot product (assigning a weight to every feature)
  - Returns a real value:  $z \in \mathbb{R}$
- How to get a probability?
  - Consider the Sigmoid function:
- Argmax for prediction:

$$\hat{y} = \arg \max_{y' \in \{0,1\}} P(y' | \mathbf{x}; \theta)$$



Logistic Regression

$$P(y | \mathbf{x}; \theta)$$

$$z = \theta \cdot \mathbf{x}$$

$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$

$$P(y = 1 | \mathbf{x}; \theta) = \sigma(\theta \cdot \mathbf{x})$$

$$P(y = 0 | \mathbf{x}; \theta) = 1 - \sigma(\theta \cdot \mathbf{x}) = \sigma(-\theta \cdot \mathbf{x})$$

# word2vec: Self-supervision

One missing piece: where to get the  $(x, y)$  pairs from?

is traditionally followed by **cherry pie**, a traditional dessert  
often mixed, such as **strawberry rhubarb pie**. Apple pie  
computer peripherals and personal **digital assistants**. These devices usually  
a computer. This includes **information available on the internet**

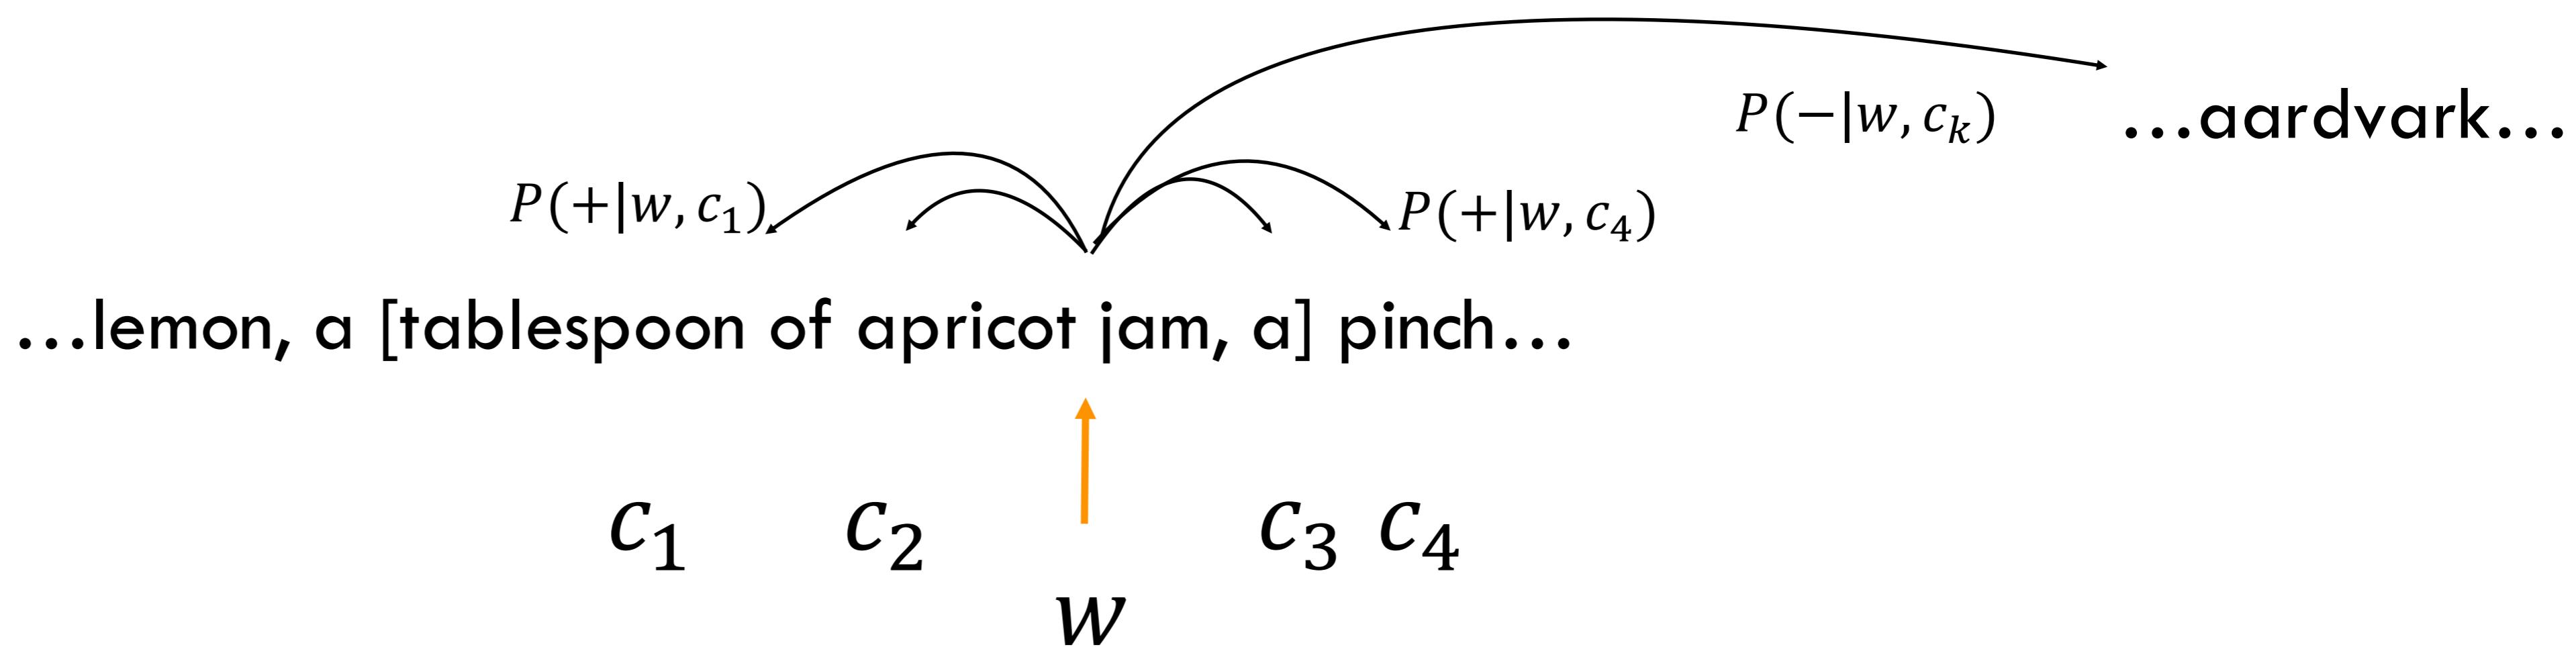
- A word  $C$  that occurs near “cherry” in the corpus acts as the gold “correct answer” for supervised learning
- No need for human labels!



What about incorrect labels?

# word2vec: Goal

Assume a +/- 2 word window, given training sentence:



Goal: train a classifier that is given a candidate (word, context) pair:

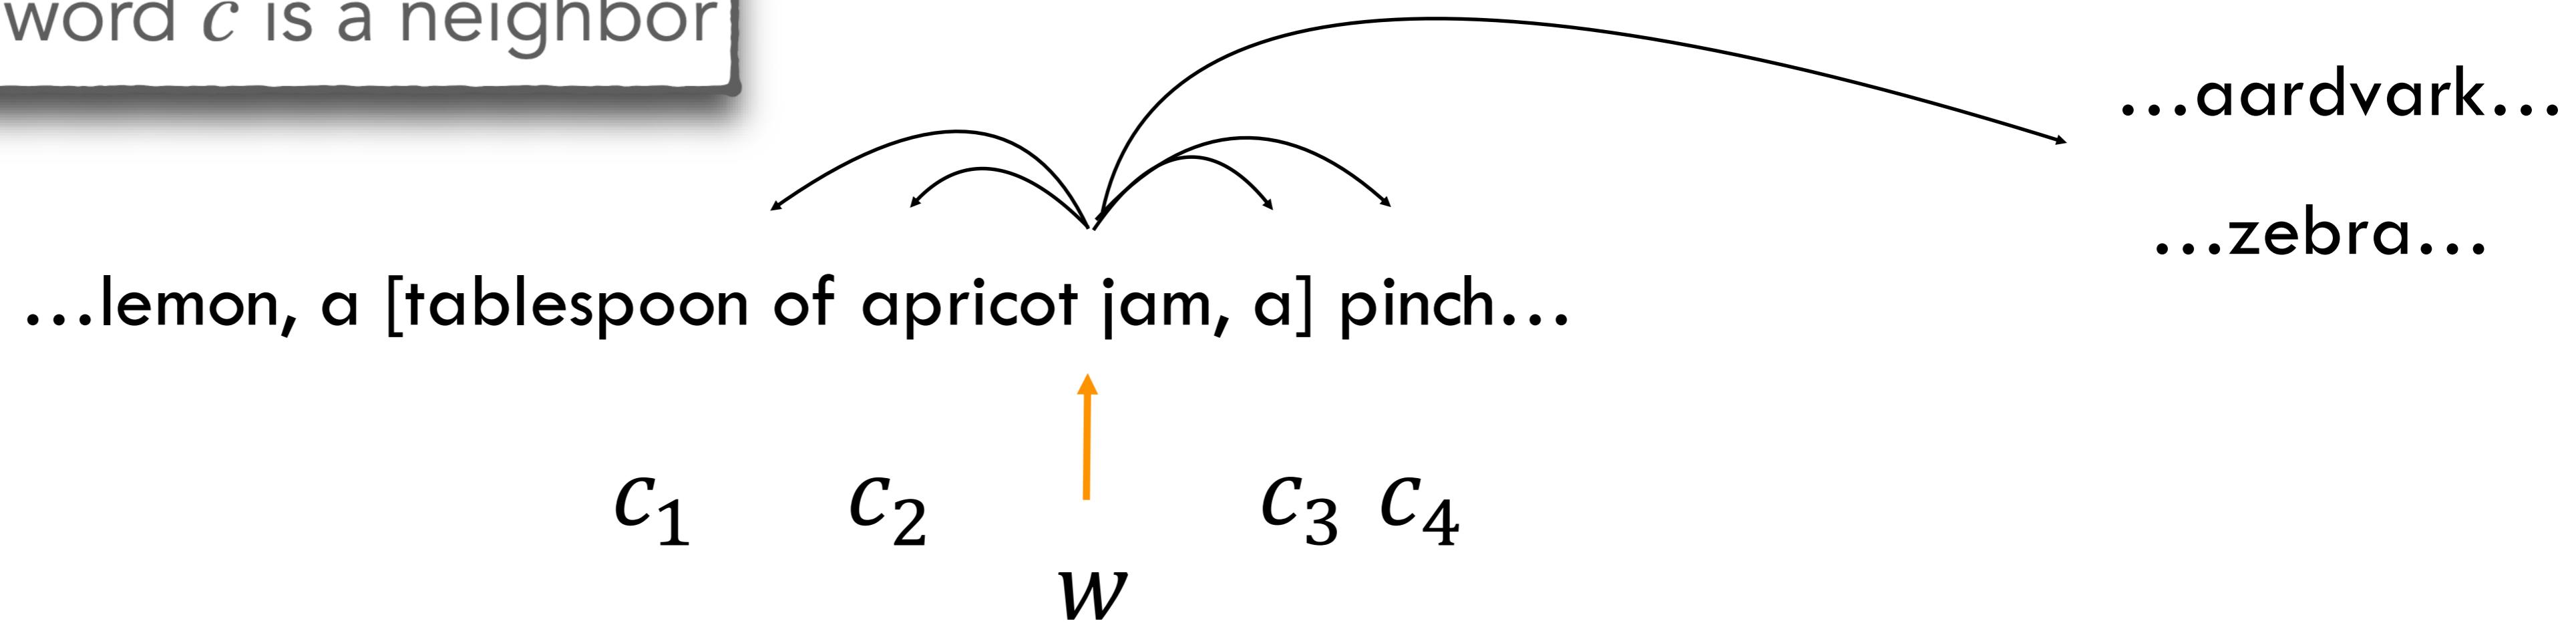
- (apricot, jam)
- (apricot, aardvark)
- ...

And assigns each pair a probability:

$$\begin{aligned} P(+|w, c) \\ P(-|w, c) = 1 - P(+|w, c) \end{aligned}$$

# word2vec: Pseudocode

Predict if candidate word  $c$  is a neighbor



1. Treat the target word  $w$  and a neighboring context word  $C$  as **positive examples**.
2. Randomly sample other words in the lexicon to get **negative examples**
3. Use logistic regression to train a classifier to distinguish those two cases
4. Use the learned weights as the embeddings